

Knowledge Representation & Processing

Yizheng Zhao^{1,2}

1. National Key Laboratory for Novel Software Technology
2. School of Artificial Intelligence, Nanjing University

Introduction to Knowledge Representation

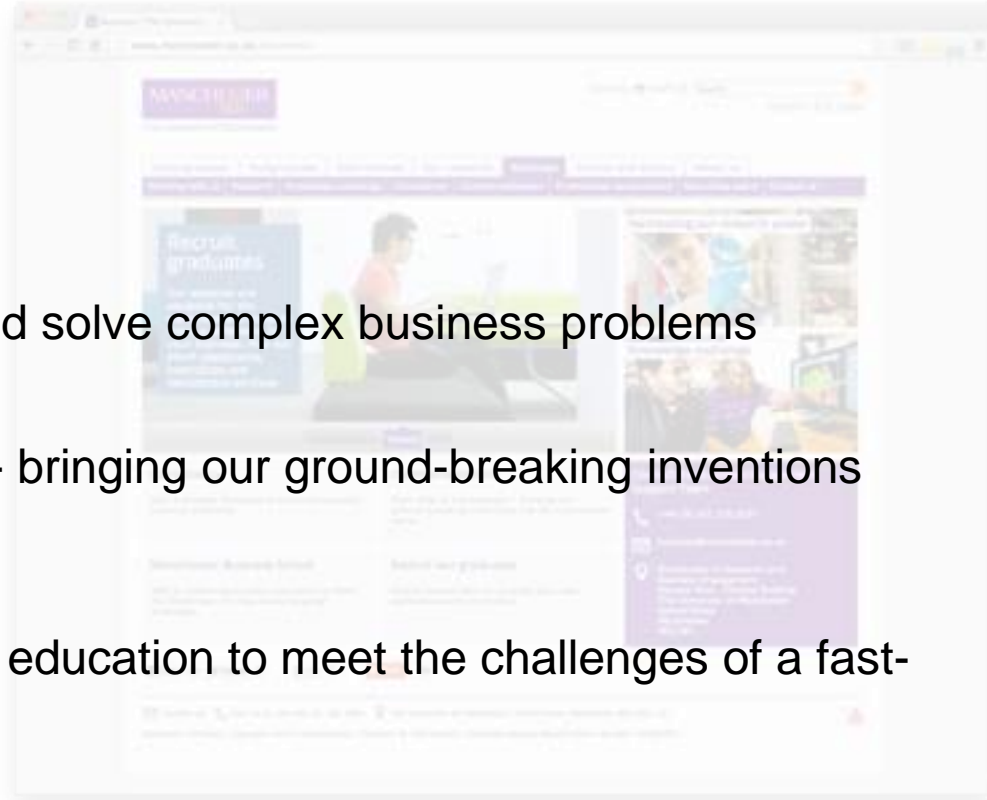
What's the Problem?




- Typical web page markup consists of:
 - Rendering information (e.g., font size and colour)
 - Hyper-links to related content
- Semantic content is accessible to humans but not (easily) to computers...

Information we can see

- University of Manchester
 - The Business School
- Consultancy
 - Gain a broader perspective and solve complex business problems
- Commercialisation
 - From idea to marketplace -- bringing our ground-breaking inventions into the commercial world
- Manchester Business School
 - MBS is redefining business education to meet the challenges of a fast-evolving global landscape
- Recruit our graduates
 - Attend careers fairs or arrange your own dedicated event on campus
- Contact the Business Engagement Support Team
 - +44 161 275 2227
 - business@manchester.ac.uk
-



Information a machine can see...



Solution: XML markup with “meaningful” tags?

[illegible]

But what about....?

<university>
</university>
<department>
</department>
<address>
</address>
<activity>
</activity>
<activity>
<details>
</details>
<activity>
<details>
</details>
<contact>
</contact>

Still the Machine only sees...

[illegible]


Need to Add “Semantics”

- External agreement on meaning of annotations
 - E.g., *Dublin Core* for annotation of library/bibliographic information
 - Agree on the meaning of a set of annotation tags
 - Problems with this approach
 - Inflexible
 - Limited
- Use Vocabulary **not** Annotations
 - Ontologies
 - New terms
 - “Conceptual Lego”
 - Meaning (**semantics**) of such terms is formally specified


Machine Processable
not
Machine Understandable

Four principles towards a Semantic Web of Data*

* With thanks to Frank van Harmelen

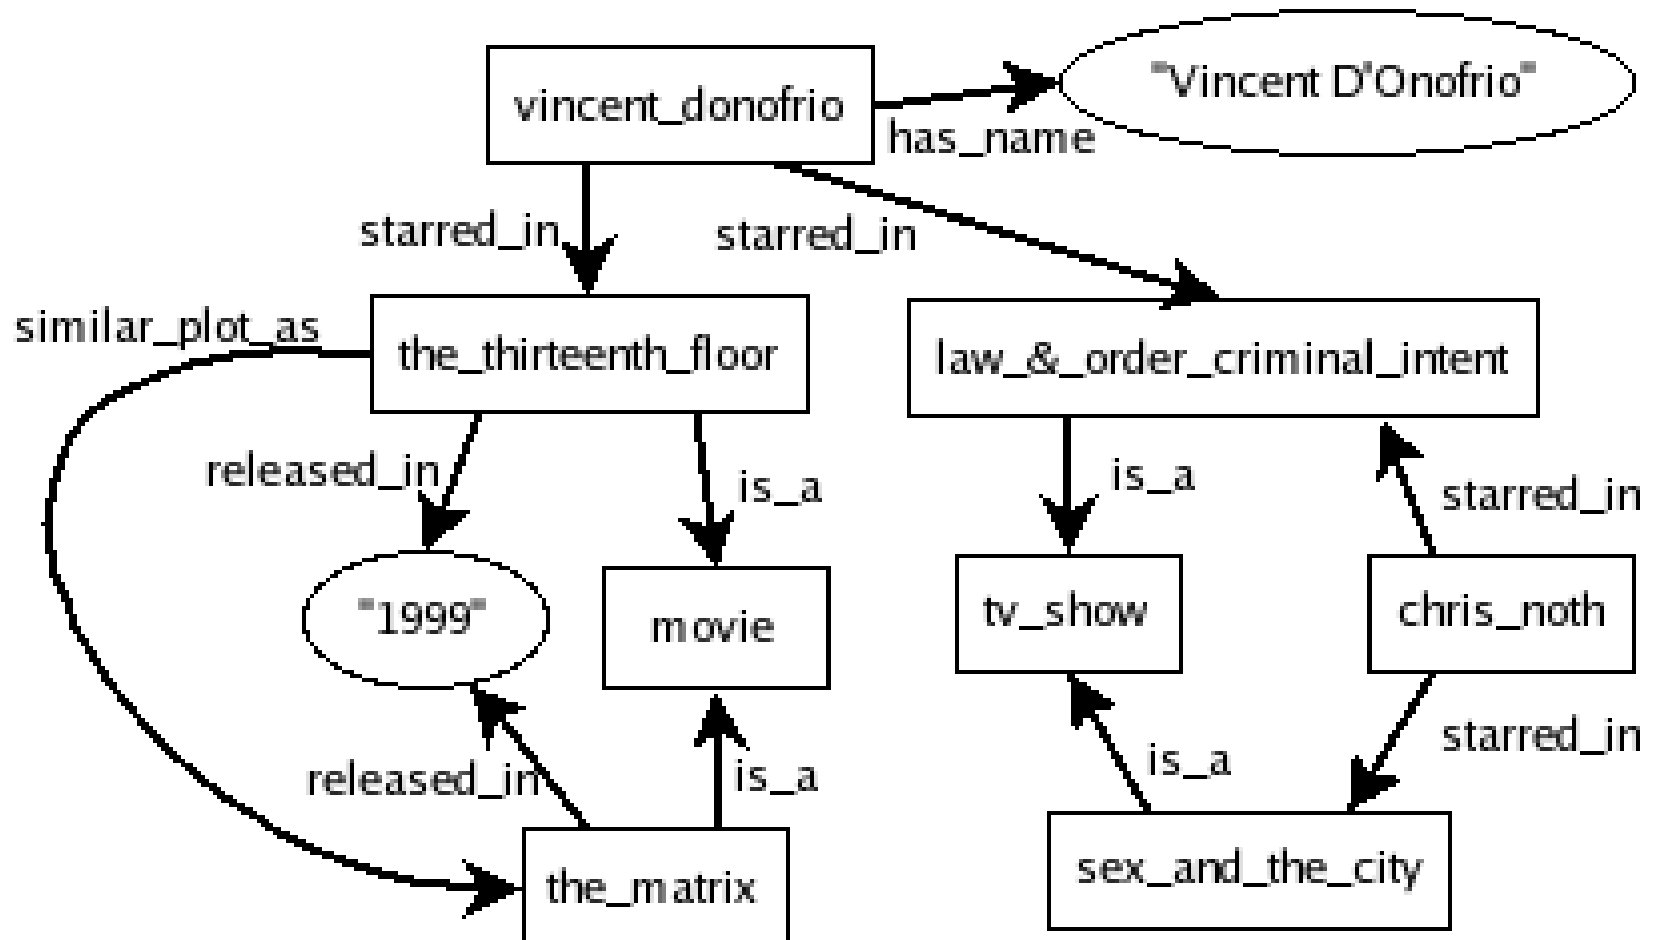


P1: Give all things a name




*“Now! That should clear up
a few things around here!”*


P2: Relationships form a graph between things



P3: The names are addresses on the Web




P1 + P2 + P3 = Giant Global Graph




P4: Explicit, Formal Semantics

- Assign Types to Things
- Assign Types to Relations
- Organise Types in a Hierarchy
- Impose Constraints on Possible Interpretations

This is where we will spend most of our time on this course unit -- looking at the ontologies that provide this semantics



Semantics



- Φρανκ is male
- married'to relates males to females

lowerbound

- married'to relates 1 male to 1 female
- Λψνδα = Ηαζελ

upperbound

Semantics = predictable inference

KR: Cloth Weaves

[Maier & Warren, Computing with Logic, 1988]

- An example showing how we can represent the qualities and characteristics of cloth types using a simple propositional logic knowledge base.




Cloth

- *Woven fabrics* consist of two sets of *threads* interlaced at right angles.
- The *warp threads* run the length of the fabric
- The *weft* (fill, pick or woof) *threads* are passed back and forth between the warp threads.
- When weaving, the warp threads are raised or lowered in patterns, leading to different weaves.
- Factors include:
 - The pattern in which warps and wefts cross
 - Relative sizes of threads
 - Relative spacing of threads
 - Colours of threads


Plain Weave

- Over and under in a regular fashion




Twill Weave

- Warp end passes over more than one weft
 - Known as “floats”
- Successive threads offset by 1



Satin Weave

- Longer “floats”
- Offsets larger than 1



Classifying Cloth

- The example provides a number of rules that describe how particular kinds of cloth are described.
- `alternatingWarp` → `plainWeave`
 - *If a piece of cloth has alternating warp, then it's a plain weave.*
- `hasFloats, warpOffsetEq1` → `twillWeave`
 - *If a piece of cloth has floats and a warp offset of 1, then it's a twill weave.*
- There are many other properties concerning the colour of threads, spacings etc.

Using the Rules

- We could use these rules to build a system that would be able to recognise different kinds of cloth through recognising the individual characteristics.
- The example given shows that once we have recognised the following characteristics
 - diagonalTexture
 - floatGTSink
 - colouredWarp
 - whiteFillthen we can determine that this cloth is denim.

Knowledge Representation

- Although this is relatively simple (in terms of both the expressivity of the language used and the number of facts), this really is an example of Knowledge Representation.
 - The rules represent some knowledge about cloth -- objects in the real world
 - Together they form a knowledge base
 - The knowledge base along with some deductive framework allow us to make inferences (which we hope reflect the characteristics/behaviour of the real world objects)

What is a Knowledge Representation?

Davis, Shrobe & Szolovits

<http://groups.csail.mit.edu/medg/ftp/psz/k-rep.html>

- Surrogate
That is, a representation
- Expression of ontological commitment
of the world
- Theory of intelligent reasoning
and our knowledge of it
- Medium of efficient computation
that is accessible to programs
- Medium of human expression
and usable

KR as Surrogate

- Reasoning is an internal process, while the things that we wish to reason about are (usually) external
- A representation acts as a surrogate, standing in for things that exist in the world.
 - Reasoning operates on the surrogate rather than the things
- Surrogates can serve for tangible and intangible objects
 - Bicycles, cats, dogs, proteins
 - Actions, processes, beliefs

KR as Surrogate

- What is the correspondence between the representation and the things it is intended to represent?
 - Semantics
- How close is the representation?
 - What's there?
 - What's missing?
- Representations are *not* completely accurate
 - Necessarily abstractions
 - Simplifying assumptions will be present
- Imperfect representation means that incorrect conclusions are inevitable.
- We can ensure that our reasoning processes are sound
 - Only guarantees that the reasoning is not the source of the error.

KR as Set of Ontological Commitments

- A representation encapsulates a collection of decisions about what to see in the world and how to see it.
- Determine the parts in focus and out of focus
 - Necessarily so because of the imperfection of representation
- Choice of representation
- Commitments as layers
- KR != Data Structure
 - Representational languages carry meaning
 - Data structures may be used to implement representations
 - Semantic Nets vs. graphs

KR as Fragmentary Theory of Intelligent Reasoning

- Incorporates only part of the insight or belief
- Insight or belief is only part of the phenomenon of intelligent reasoning
- Intelligent inference
 - Deduction
- Sanctioned inferences
 - What can be inferred
- Recommended inferences
 - What should be inferred

KR as Medium for Efficient Computation

- To use a representation, we must compute with it.
- Programs have to work with representations
 - The representation management system is a component in a larger system
 - If the representation management system is inefficient, programmers will compensate
- Representations get complex quickly
 - People need prosthetics to work well with them

KR as Medium of Human Expression

- Representations as the means by which we
 - express things about the world;
 - tell the machine about the world;
 - tell one another about the world
- Representations as a medium for communication and expression by us.
 - How general is it?
 - How precise is it?
 - Is the expressiveness adequate?
- How easy is it for us to talk or think in the representation language?
 - How easy is it? vs. can we?

KR - ontologies - OWL

- Since the conception of the *Semantic Web*, (many) people use
 - knowledge base
 - ontologysynonymously...we do here
- OWL is one language to for writing ontologies
 - just like Java is one language for writing programmes

Ontologies

- **Metadata**
 - Resources marked-up with descriptions of their content. No good unless everyone **speaks the same language**;
- **Terminologies**
 - Provide shared and common vocabularies of a domain, so search engines, agents, authors and users can communicate. No good unless everyone **means the same thing**;
- **Ontologies**
 - Provide a **shared and common understanding** of a domain that can be communicated across people and applications, and will play a major role in supporting information exchange and discovery.

Ontology

- A representation of the shared **background knowledge** for a community
- Providing the intended meaning of a formal vocabulary used to describe a certain **conceptualisation** of objects in a domain of interest
- In CS, ontology taken to mean an **engineering** artefact
- A vocabulary of terms plus **explicit characterisations** of the **assumptions** made in interpreting those terms
- Nearly always includes some notion of **hierarchical classification** (is-a)
- Richer languages allow the definition of classes through **description** of their **characteristics**
 - Introduce the possibility of using **inference** to help in management and deployment of the knowledge.


Ontologies and Ontology Representations

- “Ontology” – a word borrowed from philosophy
 - But we are necessarily building logical systems
- “Concepts” and “Ontologies”/ “conceptualisations” in their original sense are psychosocial phenomena
 - We don’t really understand them
- “Concept representations” and “Ontology representations” are engineering artefacts
 - At best approximations of our real concepts and conceptualisations (ontologies)
 - And we don’t even quite understand what we are approximating

Ontologies and Ontology Representations (cont)

- Most of the time we will just say “concept” and “ontology” but whenever anybody starts getting religious, remember...
 - *It is only a representation!*
 - We are doing engineering, not philosophy – although philosophy is an important guide
- There is no *one* way!
 - But there are consequences to different ways
 - and there are wrong ways
 - and better or worse ways for a given purposes
 - The test of an engineering artefact is whether it is fit for purpose
 - Ontology representations are engineering artefacts

A Spectrum of Representation



So why is it hard?

- Ontologies are tricky
 - People do it too easily;
People are not logicians
 - Intuitions hard to formalise
- Ontology languages are tricky
 - “All tractable languages are useless;
all useful languages are intractable”
- The evidence
 - The problem has been about for 3000 years
 - *But now it matters!*
 - *The semantic web means knowledge representation matters*

Ontology Engineering

- How do we build ontologies that are
 - Fit for purpose? (and what does that mean?)
 - Extensible?
 - Flexible?
 - Maintainable?
- Methodologies and guidelines
 - Knowledge acquisition
 - Ontology patterns
 - Normalisation
 - Upper level ontologies

Beware

- OWL is not all of Knowledge Representation
 - Knowledge Representation is not all of the Semantic Web
 - The Semantic Web is not all of Knowledge Management
 - The field is still full of controversies
-
- This course unit is to teach you about implementation in OWL